

Concurrent Programming I

HPPS

David Marchant

Based on slides by:

Troels Henriksen, Randal E. Bryant and David R. O'Hallaron


Concurrent Programming

- So far we've talked about threads and processes each running a complete programme
- Each programme has been sequential in nature
- We can spread processing over several threads or processes, each part of the same programme
- Twice the resources means twice the speed! (Nope)

The Basic Idea

- **We can break our processing over several workers**
- **Many hands make light work**

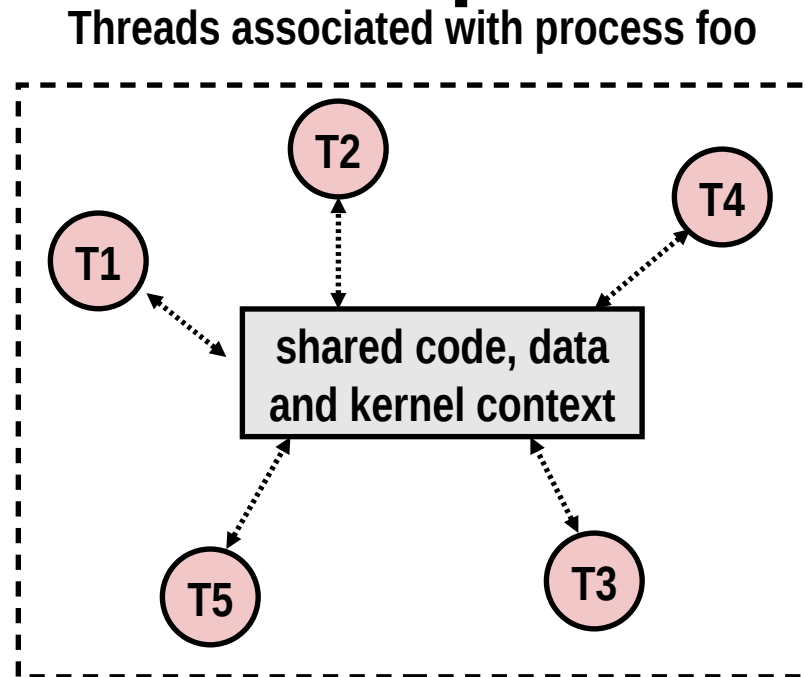
Single process
Runs for 10 mins



Two processes
Each run for 5 mins

Logical View of Threads

- Threads associated with process form a pool of peers
- Each thread can be processed in parallel




Concurrent Threads

- Two threads are *concurrent* if their flows overlap in time
- Otherwise, they are sequential

- **Examples:**


- Concurrent: A & B, A&C
- Sequential: B & C



Concurrent and Parallel Threads

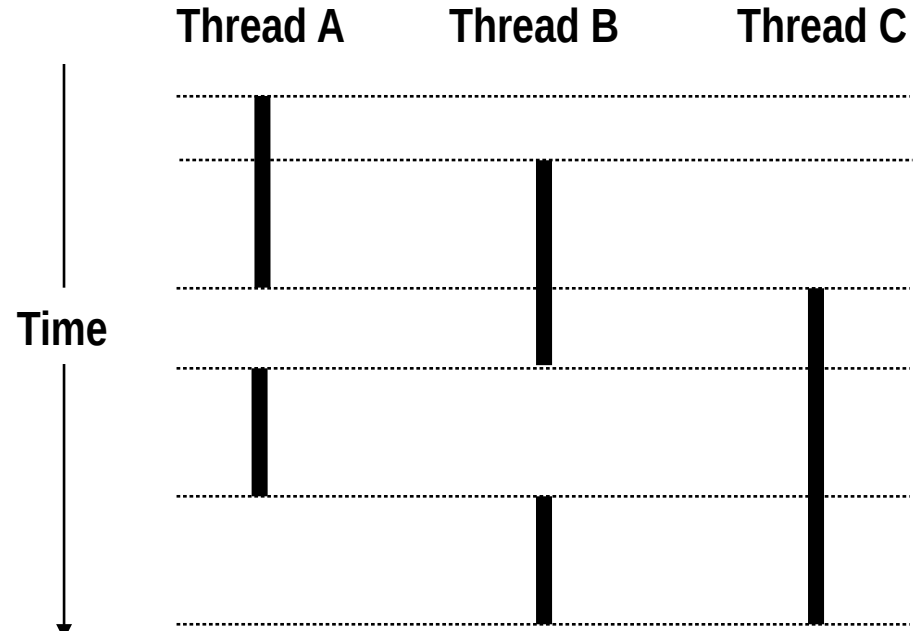
■ Single Core Processor

- Simulate parallelism by time slicing



■ Multi-Core Processor

- Can have true parallelism



Run 3 threads on 2 cores

Parallel Systems are the default

- **For many years now multi-core systems have been standard**
- **If we paid for 8 cores, we want to use them!**
- **...and not just on running background processes with slightly fewer context switches, we want a single programme to use all of the hardware**

Parallelism is about speed.

- **How do we quantify whether parallelisation made a program faster, and by how much?**
- **How can we estimate the potential benefit of parallelising a program?**
- **Are there limits to the potential gains of parallelism? So far we've talked about threads and processes each running a complete programme**

Latency

- Often called runtime
- How long it takes for a program to run on some workload.
- Both difficult and easy to measure
- *Wall time* is the time it takes in the real world.
- *CPU time* is the total amount of time spent executing code counting all processors.

Single process
Runs for 10 mins



Two processes
Each run for 5 mins

Here the CPU time is the same,
But the wall time has halved

Latency

- Suppose we have two programmes, P_1 and P_2 that each **perform the same task**, and take T_1 and T_2 respectively:

$$\text{Latency} = \frac{T_1}{T_2}$$

- Speedup greater than one means P_2 is faster than P_1 , else it is slower.
- P_1 and P_2 must solve the same problem for their latencies to be comparable.
- HINT: This is a good metric to include in any report on your own programmes

Throughput

- Sometimes latency is inappropriate—what is the running time of a potentially-infinite web server?
- Can measure latency for individual requests, but is that really useful?

$$\text{Throughput} = \frac{\text{Workload}}{\text{Time}}$$

- E.g. web requests served per second, or number of bytes processed per clock cycle.
- Using throughput we can compare programs that have different workloads.

Speedup in Throughput

- Suppose we have two programmes, P_1 and P_2 that each **perform the same task**, and have a throughput of Q_1 and Q_2 respectively:

$$\text{Throughput Speedup} = \frac{Q_1}{Q_2}$$

- Speedup greater than one means P_2 is faster than P_1 , else it is slower.
- P_1 and P_2 must solve the same problem for their throughputs to be comparable.
- HINT: This is another good metric to include in any report on your own programmes

Scalability

- **Basically how well our system improves (or not) as we make it bigger**


How the system improves in its capacity (runtime or throughput) as we add more resources, such as more processors

How the system throughput changes as we add more workload

- **Both are similar, but we are altering different parameters in our system**

Scalability Example


- **Rendering Mandelbrot fractals**
- **Each pixel can be computed independently**
- **The image size corresponds to the workload**
- **Regarding parallelism, there are two sorts of scalability that we are interested in.**



Strong Scaling

- **How the runtime varies with the number of processors for a fixed problem size.**

Speedup graph for rendering a 10^4 pixel Mandelbrot fractal as we use more threads:

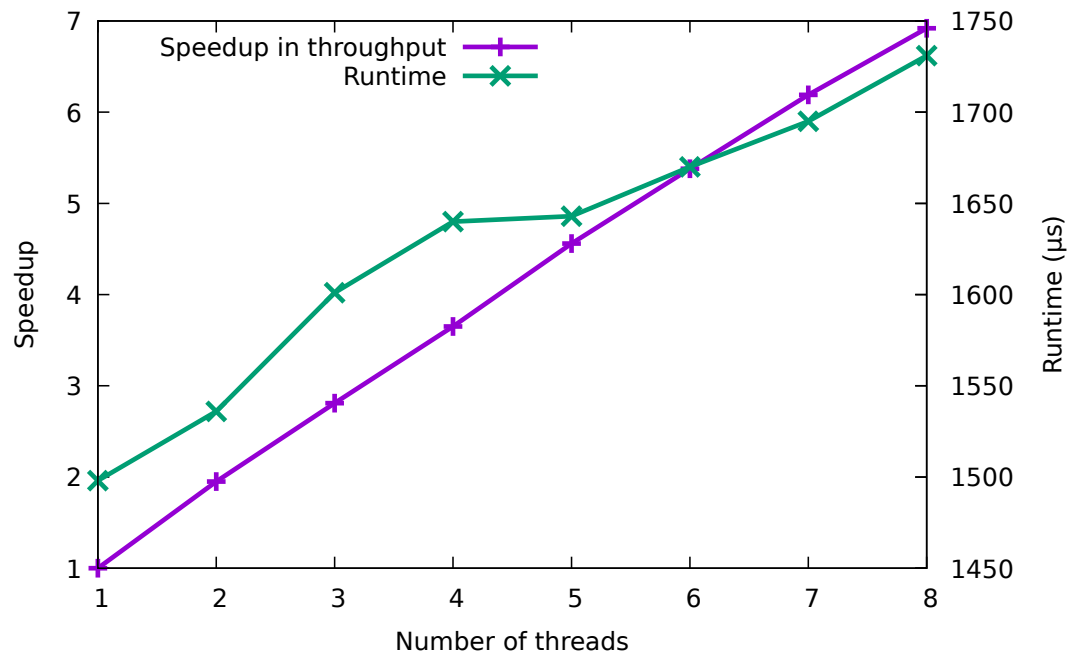


- **This shows sub-linear strong scalability—8 threads barely gets us $4 \times$ speedup**

Weak Scaling

- **How the runtime varies with the number of processors for a fixed problem size relative to the number of processors.**

Speedup graph for rendering a Mandelbrot fractal with 10^4 pixels per thread:



- **Pretty decent weak scalability! By far more common than strong scalability.**

How do we estimate the potential benefits of parallelising a program, without just implementing it and measuring how well it goes?


Amdahl's Law

- If p is the proportion of execution time that benefits from parallelisation, then $S(N)$ is maximum theoretical speedup achievable by execution on N processors, and is given by:

$$S(N) = \frac{1}{(1-p) + \frac{p}{N}}$$

- Why can p not equal 1?
- Reading, writing and other fundamentally sequential parts mean p will never be 1 in practice

Amdahl's Law



Amdahl's Law

- Amdahl's Law predicts strong scalability.
- For a fixed problem size, there will always come a point of diminishing returns.
- Makes parallelisation seem pointless...

But...

- You have a simulation that runs in 1 hour.
- Then you get a computer that is ten times faster.
- Usually you don't decide to run the simulation in six minutes, instead you make a simulation that is ten times as precise and still runs in 1 hour.
- Consider weather forecasting...
- Bigger machines let us solve bigger problems in the same time!


Gustafson's Law

- If s is the proportion of execution time that must be sequential, then $S(N)$ is maximum theoretical speedup achievable by execution on N processors, and is given by:

$$S(N) = N + (1-N) \times s$$

- Assumes that parallel workload increases just as fast as number of processors.
- Predicts weak scalability.
- In practice often more relevant than Amdahl's Law.

Gustafson's Law



The Fine Print

- **Both Amdahl's and Gustafson's Laws are idealised abstractions and ignore important real-world concerns:**
 - **Locality.**
 - **Communication.**
 - **Synchronisation.**
- **But they are still a good theoretical framework for estimating the value of parallelising some program.**
- **Neither should be used to calculate timings or parallelisation in an implemented system.**

How do we start throwing threads at our problems?

Parallel Loops

- For scientific computing, we are mostly concerned with parallelising straightforward loops such as this matrix multiplication code

```
void matmul(int n , const double *x ,  
            const double *y , double *out ) {  
    for ( int i = 0 ; i < n ; i++ ) {  
        for ( int j = 0 ; j < n ; j++ ) {  
            out [ i *n+j ] = 0 ;  
            for ( int l = 0 ; l < n ; l++ ) {  
                out [ i *n+j ] += x [ i *n+j ] * y [ i *n+j ] ;  
            }  
        }  
    }  
}
```

- Iterations of the outer two loops are independent.
- Can be computed by different threads.

Simplest OpenMP Example

```
#pragma omp parallel for  
for (int i = 0 ; i < n ; i++) {  
    A[i] = A[i] * 2;  
}
```

- ***Directives*** used to indicate how run C program in parallel.
- ***Clauses*** (covered later) can be used to customise the behaviour.
- Semantics are the ***sequential elision***—how the program would behave if we ignored the directives.
- We are only scratching the surface of OpenMP in this course!

Simplest OpenMP Example

```
print ("Program Starts\n");  
n = 1000;
```

```
#pragma omp parallel for  
for (int i=0; i<n; i++) {  
    A[i] = B[i] + C[i];  
}
```

```
m = 500;
```


```
#pragma omp parallel for  
for (int j=0; j<m; j++) {  
    p[j] = q[j] - r[j];  
}
```

```
print ("Program Ends\n");  
exit(0),
```

Sequential




Parallel



Sequential



Parallel



Sequential



- **Program starts sequential.**
- **Parallel regions split across multiple threads.**
- **Parallel region ends when all threads done.**
- **Worker threads kept running in background.**

Compilation

```
#include <stdio.h>
#include <stdlib.h>

int main (void) {
    int n = 1000000000;
    int *arr = malloc (n*sizeof(int)) ;

    #pragma omp parallel for
    for (int i=0; i<n; i++) {
        arr[ i ] = i ;
    }

    free(arr);
}
```

```
$ gcc -o openmp-example openmp-example.c -fopenmp
```


Controlling the number of threads

```
$ time OMP_NUM_THREADS=1 ./openmp-example
real    0m0.124s    Wall time
user    0m0.034s    CPU time (user)
sys     0m0.090s    CPU time (system)
$ time OMP_NUM_THREADS=2 ./openmp-example
real    0m0.076s    Wall time
user    0m0.033s    CPU time (user)
sys     0m0.104s    CPU time (system)
$ time OMP_NUM_THREADS=3 ./openmp-example
real    0m0.054s    Wall time
user    0m0.039s    CPU time (user)
sys     0m0.133s    CPU time (system)
$ time OMP_NUM_THREADS=4 ./openmp-example
real    0m0.046s    Wall time
user    0m0.054s    CPU time (user)
sys     0m0.184s    CPU time (system)
```

Memory model

- Variables declared inside a loop iteration are private.
- Variables declared outside are shared.
- Be extremely careful when modifying shared variables—OpenMP will not protect you!
- This is known as a race condition, multiple threads are reading and writing to the same location. Who gets there first? We don't know!
- We will cover this in more depth next week

```
Double sum=0;
#pragma omp parallel for
for (int i = 0 ; i < n ; i++) {
    sum += A[i];
}
```



Reductions

- Instead of doing a summation sequentially

$$(((((((x_0 + x_1) + x_2) + x_3) + x_4) + x_5) + x_6) + x_7)$$

- we can do it like:

$$(x_0 + x_1 + x_2 + x_3) + (x_4 + x_5 + x_6 + x_7)$$

- and have one compute the left part, and a second one compute the right, combining their results at the end.

Is this valid?

Reductions

- Instead of doing a summation sequentially

$$(((((((x_0 \oplus x_1) \oplus x_2) \oplus x_3) \oplus x_4) \oplus x_5) \oplus x_6) \oplus x_7)$$

- we can do it like:

$$(x_0 \oplus x_1 \oplus x_2 \oplus x_3) \oplus (x_4 \oplus x_5 \oplus x_6 \oplus x_7)$$

- and have one compute the left part, and a second one compute the right, combining their results at the end.

What about now?

Reductions

- A binary operator $\oplus : \alpha \rightarrow \alpha \rightarrow \alpha$ is said to be *commutative* if

$$x \oplus y = y \oplus x$$

- It is *associative* if

$$x \oplus (y \oplus z) = (x \oplus y) \oplus z$$

- We can do a parallel “sum” with any associative operator, called a reduction.
- For convenience, we also tend to require a neutral element 0_{\oplus} :

$$x \oplus 0_{\oplus} = 0_{\oplus} \oplus x = x$$

Reduction example

```
double dotprod (int n, double *x, double *y) {  
    double sum = 0 ;  
    #pragma omp parallel for reduction (+: sum)  
    for (int i=0; i<n; i++) {  
        sum += x[i] * y[i] ;  
    }  
    return sum ;  
}
```

- **Must initialise accumulator variable to neutral element.**
- **Must explicitly tell OpenMP the combining operator (+, *, -, &&, ||, &, |, ^, max, or min).**

Parallelising Matrix Multiplication

```
void matmul_seq(int n , const double *x ,  
                const double *y, double *out) {  
    for (int i=0; i<n; i++) {  
        for (int j=0; j<n; j++) {  
            double acc = 0;  
            for (int l=0; l<n; l++) {  
                acc += x[i*n+l] * y[j*n+l] ;  
            }  
            out[i*n+j] = acc;  
        }  
    }  
}
```

- Runtime for n=1000: 2.69s.
- Three nested loops.
- Which do we parallelise, and how?

Parallelising Matrix Multiplication

```
void matmul_seq(int n , const double *x ,
                const double *y, double *out) {
    #pragma omp parallel for
    for (int i=0; i<n; i++) {
        for (int j=0; j<n; j++) {
            double acc = 0;
            for (int l=0; l<n; l++) {
                acc += x[i*n+l] * y[j*n+l] ;
            }
            out[i*n+j] = acc;
        }
    }
}
```

- **Runtime: 0.80s (was. 2.69s seq) — pretty good for adding one line of code!**
- **But this only parallelises across the rows of the result matrix.**

Parallelising Matrix Multiplication

```
void matmul_seq(int n , const double *x ,  
                const double *y, double *out) {  
    #pragma omp parallel for collapse(2)  
    for (int i=0; i<n; i++) {  
        for (int j=0; j<n; j++) {  
            double acc = 0;  
            for (int l=0; l<n; l++) {  
                acc += x[i*n+l] * y[j*n+l] ;  
            }  
            out[i*n+j] = acc;  
        }  
    }  
}
```

- collapse(2) clause combines loops to one parallel $n*n$ iteration loop,
- Runtime: 0.93s (was 0.80s without collapse) — not much impact here.
- Would matter more if we had fewer iterations in outer loop.

Parallelising Matrix Multiplication

```
void matmul_seq(int n , const double *x ,
                const double *y, double *out) {
    for (int i=0; i<n; i++) {
        for (int j=0; j<n; j++) {
            double acc = 0;
            #pragma omp parallel for reduction(+:acc)
            for (int l=0; l<n; l++) {
                acc += x[i*n+l] * y[j*n+l] ;
            }
            out[i*n+j] = acc;
        }
    }
}
```

- Runtime: 2.48s (vs. 2.69s sequential)—this sucks.
- Almost always better to parallelise outermost loops.

Parallelising Matrix Multiplication

```
void matmul_seq(int n , const double *x ,
                const double *y, double *out) {
    #pragma omp parallel for
    for (int j=0; j<n; j++) {
        for (int i=0; i<n; i++) {
            double acc = 0;
            for (int l=0; l<n; l++) {
                acc += x[i*n+l] * y[j*n+l] ;
            }
            out[i*n+j] = acc;
        }
    }
}
```

- Runtime: 0.80s (was. 0.80 first ordering)
- No real effect here as both loop are the same, but can exploit locality for some small but easy speedup

Scheduling Clauses

- By default, OpenMP splits the iterations of a parallel loop evenly among the threads (static scheduling). This is not always optimal.

```
int fib (int n) {  
    if (n <= 1) {  
        return 1;  
    } else {  
        return fib(n-1) + fib(n-2);  
    }  
}
```

- Consider a loop that computes fib(i) for each $i < n$.
- Since time to compute fib(i) is over twice that of fib(i-1), the threads with the early iterations finish much faster than the threads with the later iterations.

```
#pragma omp parallel for schedule(static)  
for (int i=0; i<n; i++) {  
    fibs[i] = fib(i);  
}
```


Dynamic Scheduling

- Dynamic scheduling assigns each thread an iteration, and is given more iterations when it finishes.
- No idle threads (as long as there are unclaimed iterations to run).

```
#pragma omp parallel for schedule(dynamic)  
for (int i=0; i<n; i++) {  
    fibs[i] = fib(i);  
}
```

- On my machine static ran in 5.2s, where dynamic ran in 2.27s—much better!
- By default, assigns single iterations at a time, which is slow for very large loops with quick iterations.
- Use `schedule(dynamic,K)` to schedule K iterations at a time to threads.

Summary

- **Parallelism is about speedup.**
- **Describe performance differences with speedup.**
 - **Latency for programs that compute the same thing.**
 - **Throughput if not.**
- **Strong scaling is solving the same problem faster.**
- **Weak scaling is solving a bigger problem in the same time.**
- **OpenMP is a simple language extension for parallelising loops in C programs.**
- **Use `#pragma omp parallel for` to parallelise a for-loop.**
- **Generally parallelising the outer most loop works best**